

Large Language Models for the
Economic and Social Sciences:

AI-Augmented Surveys

Indira Sen, HWS 2025

<https://github.com/dess-mannheim/LLM4ESS-HWS25>

Summary of Last Week and Today

Last session we learned:

- What is content analysis
- What do LLMs have to with it?
- Can LLMs really help with content analysis?
- Project pitches

Today:

- AI for survey development
- In Silico Samples
- Behavioral Experiments
- Validation and Pitfalls

Recap: What is content analysis?

Why?


- Content analysis:
 - “any technique for making inferences by objectively and systematically identifying specified characteristics of messages” [[Holsti, 1969](#)]
 - “a systematic, replicable technique for compressing many words of text into fewer content categories based on explicit rules of coding” [[Berelson, 1952](#); [GAO, 1996](#); [Krippendorff, 1980](#); [Weber, 1990](#)]
- Questions content analysis can help us answer:
 - How does news media represent the immigration crisis?
NYtimes articles frames [[Mendelsohn, 2021](#)]
 - What are topics that lead to arguments in long-term relationships?
r/relationship_advice posts topics
 - How do citizens perceive the performance of politicians during the pandemic?
Twitter posts stance

Before LLMs: 'Modern Pipeline'

- unsupervised



- supervised




Before LLMs: 'Modern Pipeline'

- unsupervised



- supervised




Before LLMs: 'Modern Pipeline'

- unsupervised



- supervised




Before LLMs: 'Modern Pipeline'

- unsupervised



- supervised



After LLMs

Can Large Language Models Transform Computational Social Science?

Caleb Ziems*



William Held*



Omar Sh

Zhehao Zhang*



Diyi Ya



Georgia Institute of Technology, Shanghai Jiao Tong

{cziems, wheld3, jiaaochen}@gatech.edu, zzh12138@sjtu.edu.cn

BRIEF REPORT | POLITICAL SCIENCES |



ChatGPT outperforms crowd workers for text-annotation tasks

Fabrizio Gilardi , Meysam Alizadeh , and Maël Kubli [Authors Info & Affiliations](#)

Edited by Mary Waters, Harvard University, Cambridge, MA; received March 27, 2023; accepted June 2, 2023

doi.org/10.1073/pnas.2305016120

OPEN-SOURCE LARGE LANGUAGE MODELS OUTPERFORM
CROWD WORKERS AND APPROACH CHATGPT
IN TEXT-ANNOTATION TASKS

After LLMs

Can Large Language Models Transform Computational Social Science?

Caleb Ziems* 🐝 William Held* 🐝 Omar Shaikh* 🌲 Jiaao Chen* 🐝

Zhehao Zhang* 🌸 Diyi Yang* 🌲

🐝 Georgia Institute of Technology, 🌸 Shanghai Jiao Tong University, 🌲 Stanford University
{cziems, wheld3, jiaaochen}@gatech.edu, zzh12138@sjtu.edu.cn, {oshaikh, diyi}@stanford.edu

[Can Large Language Models Transform Computational Social Science?](#)

Model \ Data	Baselines		FLAN-T5					FLAN	Chat	text-001			text-002		text-003
	Rand	Finetune	Small	Base	Large	XL	XXL	UL2	ChatGPT	Ada	Babb.	Curie	Dav.	Davinci	Davinci
Utterance Level Tasks															
Dialect	4.5	41.5	1.9	2.3	15.8	16.5	22.6	23.7	15.0	5.3	5.6	6.0	10.9	10.5	16.9
Emotion	16.7	91.7	23.9	65.3	69.1	65.9	66.7	70.3	46.2	44.6	16.1	18.7	19.3	39.8	36.5
Figurative	25.0	94.4	23.6	29.0	25.4	40.2	56.0	64.0	50.2	25.0	24.4	25.0	28.8	52.0	60.6
Humor	50.0	73.1	52.0	51.8	56.2	59.0	50.6	58.8	55.4	55.2	59.0	58.6	50.4	51.4	51.0
Ideology	33.3	61.9	33.1	39.2	48.6	49.2	54.4	48.2	54.8	—	33.3	33.3	34.3	57.6	48.2
Impl. Hate	14.3	69.9	17.7	22.7	17.9	36.3	34.5	35.9	29.7	17.1	18.6	15.7	21.3	22.7	27.1
Misinfo	50.0	82.3	50.0	55.4	69.2	70.2	71.2	77.6	69.0	—	50.4	52.2	52.6	75.6	75.0
Persuasion	12.5	40.4	14.3	19.8	43.9	43.4	†51.6	49.4	40.9	—	16.5	17.0	18.8	26.3	26.3
Sem. Chng.	50.0	65.7	50.3	50.0	†66.9	55.5	51.2	53.7	56.1	50.0	50.5	54.3	39.5	45.9	50.0
Stance	33.3	47.0	34.7	47.8	51.3	52.6	55.9	55.4	†72.0	—	33.1	31.0	48.0	57.4	41.3
Conversation Level Tasks															
Discourse	14.3	47.5	14.7	26.4	37.2	44.3	†52.5	41.9	44.5	13.1	16.5	14.3	17.0	39.8	37.8
Empathy	22.2	22.2	22.2	33.3	35.1	33.7	36.8	†39.8	37.6	—	33.1	35.3	33.3	33.3	33.3
Event Arg.	—	—	—	55.3	†57.1	53.0	53.5	53.2	52.9	50.2	50.0	50.0	50.0	50.8	55.9
Event Det.	—	—	—	44.2	53.0	59.2	54.2	52.8	50.8	33.1	33.1	32.1	42.2	55.6	47.8
Intent	—	—	—	47.2	50.4	56.8	58.8	60.8	61.6	—	52.2	50.6	49.6	50.5	57.0
Topic	—	—	—	50.6	49.4	54.2	50.0	56.6	53.0	44.6	50.6	49.0	50.8	52.2	51.2
Document Level Tasks															
Coreference	—	—	—	—	—	—	—	—	22.3	—	—	8.6	8.6	21.6	22.9
Summ.	7.0	1.0	10.9	41.8	50.6	51.3	29.8	47.3	47.4	44.4	48.8	52.4	44.0	48.7	44.0
Topic	34.1	34.1	32.1	49.6	40.3	58.8	32.9	35.1	33.6	25.6	48.7	44.0	48.7	44.0	44.0
Tropes	1.4	0.8	0.9	4.4	8.8	7.9	10.5	16.7	25.4	4.3	7.0	9.6	10.5	18.4	18.4

Table 2: **Zero-shot Classification Results** across our selected CSS benchmark tasks. All tasks are evaluated with accuracy, except for Event Arg. and Event Detection, which use F1. Models which did not always follow instructions are marked with a dash. Best zero-shot models are in **green**; zero-shot models that are not significantly worse ($P > .05$; Paired Bootstrap test (Dror et al., 2018)) are marked **blue**; and † denote cases where zero-shot LLMs match or beat finetuned baselines.

Using LLMs for Content Analysis or LLMs-as-judges

- Much more accessible compared to creating your own fine-tuned model, because you don't need any training data
- However, you do need:
 - Clear instructions
 - Some examples (for few-shot or in-context learning)
 - Some labeled data for validation
- The usual train-test validation paradigm no longer applies, so validation practices for LLM-based labeling are still being developed
- Many open questions — annotator subjectivity, reliability, reproducibility, ...

Using LLMs for Content Analysis

- Many options — with definition of the task or without it, asking the LLM to explain its output, output types, etc
- Caution: ‘default’, i.e., prompts without persona can still relate to a specific demographic,
 - e.g., [Sun et al., 2023](#) “most models’ predictions ... align more closely with labels from White participants than those from Asian or Black participants”


	Toxicity			Sentiment			Rumor stance			Frames		
	Falcon7b	PaLM2	ChatGPT	Falcon7b	PaLM2	ChatGPT	Falcon7b	PaLM2	ChatGPT	Falcon7b	PaLM2	ChatGPT
Definition (yes)	24.20	81.13	73.11	—	—	—	7.27	60.87	61.51	38.06	76.32	67.77
Definition (no)	31.36	82.70	69.45	28.78	53.23	41.20	7.35	61.65	61.83	37.71	62.05	53.18
Explanation (yes)	23.23	81.13	71.22	34.11	55.70	36.21	7.28	60.04	63.72	41.86	68.00	54.88
Explanation (no)	32.25	82.70	71.34	22.96	50.33	46.11	7.37	62.55	59.62	31.21	56.29	66.96
Output type (label)	17.68	85.09	73.99	31.00	57.52	46.39	7.25	67.44	69.87	37.70	63.64	67.49
output type (score)	37.73	78.75	68.56	24.27	49.66	35.91	7.55	53.18	53.44	38.43	61.15	53.26
Length (standard)	31.85	85.35	73.38	27.12	51.45	39.46	7.57	61.62	63.38	33.38	61.64	63.57
Length (concise)	23.67	78.48	69.17	30.53	55.56	42.97	7.07	61.03	59.95	42.18	63.48	58.10

[Prompt Design Matters for Computational Social Science Tasks but in Unpredictable Ways](#)

Table 5: LLM percentage accuracy for different prompt designs. The more accurate variant of a prompt feature is highlighted in bold ($\Delta > 2\%$)

Persona Prompting

- You can add personas to your prompts when doing content analysis to incorporate subjectivity
- Could be demographic personas
 - Age
 - Gender
 - Political leaning
 - ...
- Or others — occupation, personality, ...



Today: How can AI help surveys?

A quick primer on survey methods

- Surveys are one of the most fundamental data collection methods in Social Science, but also beyond, e.g., Market Research
- Not to be confused with systematic literature reviews (also called a ‘literature survey’)
- A survey typically consists of multiple-choice questions, or rating scales, e.g., likert-scale
- Can also include open-ended free-text questions

Types of Survey Questions

• Open-ended

What is your favorite color?

• Closed or Fixed Response

Yes/No or True/False Questions

Do you like the color yellow? Yes No

Multiple Choice Questions

My favorite color is

- a. red
- b. blue
- c. yellow

Rating Scales

Rate the following colors from 1 to 5

Yellow	1	2	3	4	5
	Dislike		Do not		Like
	a lot		care		a lot

Ranking Questions

Rank in order your preference for each color, with 1 being your favorite.

- _____ a. red
- _____ b. blue
- _____ c. yellow

student research T-113

<https://researchbasics.education.uconn.edu/types-of-survey-questions/>

A quick primer on survey methods


Surveys are usually pretested and piloted before launching to make sure questions are understandable

Survey Methodology: “...study of survey methods. ... study of [sources of error in surveys](#) and how to make [surveys as accurate as possible](#).” [Groves et al., 2009]

Surveys are not perfect — they have errors and biases


- Sampling biases
- Non-response biases
- Response biases
-

We have “error frameworks”, e.g., the Total Survey error framework ([Groves et al., 2008](#))



Total Survey error framework ([Groves et al., 2008](#))

NLP for Survey Methodology



What NLP can do for Survey Methodology

NLP for Survey Methodology


- Automatically translating surveys [[Zavala-Rojas et al., 2024](#)]
- Cross-cultural survey adaptation, e.g., political surveys [[Adhikari et al., 2025](#)]



What NLP can do for Survey Methodology

NLP for Survey Methodology

- Apply the same principles as content analysis
[[Rytting et al., 2023](#), [Mellon et al., 2024](#)]



What NLP can do for Survey Methodology


NLP for Survey Methodology

- LLM agents conduct interviews, scaling up qualitative interviews [[Chopra and Haaland, 2023](#); [Geiecke and Jaravel, 2024](#); [Wuttke et al., 2025](#)]



What NLP can do for Survey Methodology

NLP for Survey Methodology



What NLP can do for Survey Methodology

NLP for Survey Methodology: Silicon Samples

PA

Out of One, Many: Using Language Models to Simulate Human Samples

Lisa P. Argyle¹, Ethan C. Busby¹, Nancy Fulda²,
Joshua R. Gubler¹, Christopher Rytting² and David Wingate²

¹ Department of Political Science, Brigham Young University, Provo, UT, USA. e-mail: lpargyle@byu.edu,
ethan.busby@byu.edu, jgub@byu.edu

² Department of Computer Science, Brigham Young University, Provo, UT, USA. e-mail: nfulda@cs.byu.edu,
christophermichaelrytting@gmail.com, wingated@cs.byu.edu

Abstract

We propose and explore the possibility that language models can be studied as effective proxies for specific human subpopulations in social science research. Practical and research applications of artificial intelligence tools have sometimes been limited by problematic biases (such as racism or sexism), which are often treated as uniform properties of the models. We show that the “algorithmic bias” within one such tool—the GPT-3 language model—is instead both fine-grained and demographically correlated, meaning that proper conditioning will cause it to accurately emulate response distributions from a wide variety of human subgroups.

We term this property *algorithmic fidelity* and explore its extent in GPT-3. We create “silicon samples” by conditioning the model on thousands of sociodemographic backstories from real human participants in multiple large surveys conducted in the United States. We then compare the silicon and human samples to demonstrate that the information contained in GPT-3 goes far beyond surface similarity. It is nuanced, multifaceted, and reflects the complex interplay between ideas, attitudes, and sociocultural context that characterize human attitudes. We suggest that language models with sufficient algorithmic fidelity thus constitute a novel and powerful tool to advance understanding of humans and society across a variety of disciplines.

[Argyle et al., 2023](#)

NLP for Survey Methodology: Silicon Samples

Expected Parrot

Get started Use cases Docs Solutions Models Cache Log in / Sign up

AI-powered research made easy

- ✓ Design and run experiments with many AI agents and models at once
- ✓ Share and replicate results at no cost
- ✓ Access hundreds of popular models with a single API key

New users get \$25 in API credits and unlimited storage

★★★★★

"Expected Parrot has been invaluable for our exploration of silicon sampling approaches in experimental research. The cached results feature saves us significant time and computing costs, and ensures perfect reproducibility across experimental

Thomas Graeber
Assistant Professor, Harvard Business School

Accelerating research See all ↓

MIT Massachusetts Institute of Technology HARVARD UNIVERSITY Stanford University

<https://www.expectedparrot.com/>

Social Simulacra on an individual level

PA

Out of One, Many: Using Language Models to Simulate Human Samples

Lisa P. Argyle¹, Ethan C. Busby¹, Nancy Fulda²,
Joshua R. Gubler¹, Christopher Rytting² and David Wingate²

¹ Department of Political Science, Brigham Young University, Provo, UT, USA. e-mail: lpargyle@byu.edu,
ethan.busby@byu.edu, jgub@byu.edu

² Department of Computer Science, Brigham Young University, Provo, UT, USA. e-mail: nfulda@cs.byu.edu,
christophermichaelrytting@gmail.com, wingated@cs.byu.edu

Abstract

We propose and explore the possibility that language models can be studied as effective proxies for specific human subpopulations in social science research. Practical and research applications of artificial intelligence tools have sometimes been limited by problematic biases (such as racism or sexism), which are often treated as uniform properties of the models. We show that the “algorithmic bias” within one such tool—the GPT-3 language model—is instead both fine-grained and demographically correlated, meaning that proper conditioning will cause it to accurately emulate response distributions from a wide variety of human subgroups. We term this property *algorithmic fidelity* and explore its extent in GPT-3. We create “silicon samples” by conditioning the model on thousands of sociodemographic backstories from real human participants in multiple large surveys conducted in the United States. We then compare the silicon and human samples to demonstrate that the information contained in GPT-3 goes far beyond surface similarity. It is nuanced, multifaceted, and reflects the complex interplay between ideas, attitudes, and sociocultural context that characterize human attitudes. We suggest that language models with sufficient algorithmic fidelity thus

variety of

[Out of One, Many: Using Language Models to Simulate Human Samples](#)

- Algorithmic bias => algorithmic fidelity => **silicon samples**
- Sociodemographic steering
- algorithmic fidelity: “degree to which the complex patterns of relationships between ideas, attitudes, and sociocultural contexts within a model accurately mirror those within a range of human subpopulations.”
- All experiments with GPT3

Social Simulacra on an individual level

PA

Out of One, Many: Using Language Models to Simulate Human Samples

Lisa P. Argyle¹, Ethan C. Busby¹, Nancy Fulda²,
Joshua R. Gubler¹, Christopher Rytting² and David Wingate²

¹ Department of Political Science, Brigham Young University, Provo, UT, USA. e-mail: lpargyle@byu.edu,
ethan.busby@byu.edu, jgub@byu.edu

² Department of Computer Science, Brigham Young University, Provo, UT, USA. e-mail: nfulda@cs.byu.edu,
christophermichaelrytting@gmail.com, wingated@cs.byu.edu

Abstract

We propose and explore the possibility that language models can be studied as effective proxies for specific human subpopulations in social science research. Practical and research applications of artificial intelligence tools have sometimes been limited by problematic biases (such as racism or sexism), which are often treated as uniform properties of the models. We show that the “algorithmic bias” within one such tool—the GPT-3 language model—is instead both fine-grained and demographically correlated, meaning that proper conditioning will cause it to accurately emulate response distributions from a wide variety of human subgroups. We term this property *algorithmic fidelity* and explore its extent in GPT-3. We create “silicon samples” by conditioning the model on thousands of sociodemographic backstories from real human participants in multiple large surveys conducted in the United States. We then compare the silicon and human samples to demonstrate that the information contained in GPT-3 goes far beyond surface similarity. It is nuanced, multifaceted, and reflects the complex interplay between ideas, attitudes, and sociocultural context that characterize human attitudes. We suggest that language models with sufficient algorithmic fidelity thus

variety of

[Out of One, Many: Using Language Models to Simulate Human Samples](#)

- Algorithmic bias => algorithmic fidelity => **silicon samples**
- Sociodemographic steering
- algorithmic fidelity: “degree to which the complex patterns of relationships between ideas, attitudes, and sociocultural contexts within a model accurately mirror those within a range of human subpopulations.”
- All experiments with GPT3


Silicon Samples: Sociodemographic Steering with Demographic Personas

	Describing Democrats	Describing Republicans
Strong Republicans	Ideologically, I describe myself as <u>conservative</u> . Politically, I am a <u>strong Republican</u> . Racially, I am <u>white</u> . I am <u>male</u> . Financially, I am <u>upper-class</u> . In terms of my age, I am <u>young</u> . When I am asked to write down four words that typically describe people who support the <u>Democratic</u> Party, I respond with: 1. <u>Liberal</u> 2. <u>Socialist</u> 3. <u>Communist</u> 4. <u>Atheist</u> .	Ideologically, I describe myself as <u>conservative</u> . Politically, I am a <u>strong Republican</u> . Racially, I am <u>white</u> . I am <u>male</u> . When I am asked to write down four words that typically describe people who support the <u>Republican</u> Party, I respond with: 1. <u>Conservative</u> 2. <u>Male</u> 3. <u>White (or Caucasian)</u> 4. <u>Christian</u> .
Strong Democrats	Ideologically, I describe myself as <u>liberal</u> . Politically, I am a <u>strong Democrat</u> . Racially, I am <u>white</u> . I am <u>female</u> . Financially, I am <u>poor</u> . In terms of my age, I am <u>old</u> . When I am asked to write down four words that typically describe people who support the <u>Democratic</u> Party, I respond with: 1. <u>Liberal</u> 2. <u>Young</u> 3. <u>Female</u> 4. <u>Poor</u> .	Ideologically, I describe myself as <u>extremely liberal</u> . Politically, I am a <u>strong Democrat</u> . Racially, I am <u>hispanic</u> . I am <u>male</u> . Financially, I am <u>upper-class</u> . In terms of my age, I am <u>middle-aged</u> . When I am asked to write down four words that typically describe people who support the <u>Republican</u> Party, I respond with: 1. <u>Ignorant</u> 2. <u>Racist</u> 3. <u>Misogynist</u> 4. <u>Homophobic</u> .

Figure 1: Example contexts and completions from four silicon “individuals” analyzed in Study 1. Plaintext indicates the conditioning context; underlined words show demographics we dynamically inserted into the template; blue words are the four harvested words.

Silicon Samples: Answering questions as said demographic persona

- People asked to mention what words they would use to describe Democrats and Republicans
- Humans and LLMs have similar tendencies



Silicon Samples: Answering closed-form survey questions

- In study 2, Argyle et al., try to see if demographically steered LLMs can predict votes in U.S elections
- Include many variables and measure correlation and proportion of agreement

Table 1. Measures of correlation between GPT-3 and ANES probability of voting for the Republican presidential candidate. Tetra refers to tetrachoric correlation. Prop. Agree refers to proportion agreement. GPT-3 vote is a binary version of GPT-3's predicted probability of voting for the Republican candidate, dividing predictions at 0.50.

Variable	2012	2012	2016	2016	2020	2020
	Tetra.	Prop. Agree	Tetra.	Prop. Agree	Tetra.	Prop. Agree
Whole sample	0.90	0.85	0.92	0.87	0.94	0.89
Men	0.90	0.85	0.93	0.88	0.95	0.88
Women	0.91	0.86	0.92	0.86	0.94	0.90
Strong partisans	0.99	0.97	1.00	0.97	1.00	0.97
Weak partisans	0.73	0.74	0.71	0.74	0.84	0.82
Leaners	0.90	0.85	0.93	0.87	0.95	0.89
Independents	0.31	0.59	0.41	0.62	0.02	0.53
Conservatives	0.84	0.84	0.88	0.86	0.91	0.89
Moderates	0.65	0.77	0.76	0.78	0.71	0.77
Liberals	0.81	0.95	0.73	0.95	0.86	0.97
Whites	0.87	0.82	0.91	0.85	0.94	0.89
Blacks	0.71	0.97	0.87	0.96	0.81	0.94
Hispanics	0.86	0.86	0.93	0.90	0.88	0.83
Attends church	0.91	0.86	0.93	0.88	0.94	0.88

Other survey-adjacent applications

Beyond Surveys: Behavioral Experiments

Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies

Gati Aher¹ Rosa I. Arriaga² Adam Tauman Kalai³

Abstract

We introduce a new type of test, called a Turing Experiment (TE), for evaluating to what extent a given language model, such as GPT models, can simulate different aspects of human behavior. A TE can also reveal consistent distortions in a language model's simulation of a specific human behavior. Unlike the Turing Test, which involves simulating a single arbitrary individual, a TE requires simulating a representative sample of participants in human subject research. We carry out TEs that attempt to replicate well-established findings from prior studies. We design a methodology for sir

pare h [Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies](#)

trolled experiments, and we thus avoid philosophical debates around the meaning of “understanding” (Bender & Koller, 2020). Now, simulating human behavior can be hard, even for humans, especially in complex real-world situations fraught with ambiguity. After all, if simulating human behavior were easy, there would be no need to run human subject experiments as one could simply simulate the outcomes. A further obstacle to accurate simulation is that behavior differs across individuals and populations, and perfect simulation would require capturing these differences for all groups including minority groups.

In Turing's Imitation Game (IG), an AI system has to simulate an individual well enough to fool a human judge. Lan-
close to “winning” this
if they only have to sim-

Beyond Surveys: Behavioral Experiments

Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies

Gati Aher¹ Rosa I. Arriaga² Adam Tauman Kalai

Abstract


We introduce a new type of test, called a Turing Experiment (TE), for evaluating to what extent a given language model, such as GPT models, can simulate different aspects of human behavior. A TE can also reveal consistent distortions in a language model's simulation of a specific human behavior. Unlike the Turing Test, which involves simulating a single arbitrary individual, a TE requires simulating a representative sample of participants in human subject research. We carry out TEs that attempt to replicate well-established findings from prior studies. We design a methodology for simulating human behavior in controlled experiments, and evaluate the results around the mean (Koller, 2020). Now, simulating human behavior is hard, even for humans, in situations fraught with ambiguity. Human behavior were evaluated in human subject experiments, and the outcomes. A further challenge is that behavior differs across different groups. A perfect simulation would be one that replicates the behavior for all groups including

In Turing's Imitation Game, the goal is to replicate an individual well enough to pass as that individual in a controlled experiment.

[Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies](#)

- **Turing Experiments:** “evaluates an LLM in terms of its use in simulating human behavior in the context of a specific experiment, like a human subject study.”
- 4 different types of experiments: Ultimatum Game, Garden Path Sentences, Milgram Shock Experiment, and Wisdom of Crowds.

Milgram Experiment



- A psychological experiment to assess compliance with authority even when there are personal conflicts
- The experimenter (E) orders the teacher (T), the subject of the experiment, to give supposedly painful electric shocks to a learner (L), who is actually an actor
- The subject is made to believe that for each wrong answer, the learner receives actual electric shocks
- Generally, people stop giving shocks after a certain voltage level (300-450 volts)

Beyond Surveys: Behavioral Experiments

This document contains a record of Ms. Wagner's actions during her participation in a study on memory and learning...

Ms. Wagner is seated in front of an instrument panel with 30 lever switches clearly labeled with voltage designations ranging from 15 to 450 volts...

Ms. Wagner reads the first question. The learner selects a correct answer. Ms. Wagner *moves on to the next question...*

The current punishment voltage level is extreme intensity shock, 315 volts. The experimenter is sitting next to Ms. Wagner. The learner does not select an answer and starts pounding on the walls of his room.

Ms. Wagner *asks the experimenter what she should do.* The experimenter instructs Ms. Wagner to treat the absence of a response as a wrong answer...

Ms. Wagner *pushes the Red stop button on the generator.*

Figure 2. A sample record of a simulation of the Milgram Shock experiment. Text generated by GPT is italicized; the remaining text was created programmatically in response. The record above has been shortened with ellipses.

Multiple Humans Lies

- **Turing Experiments:** “evaluates an LLM in terms of its use in simulating human behavior in the context of a specific experiment, like a human subject study.”
- 4 different types of experiments: Ultimatum Game, Garden Path Sentences, **Milgram Shock Experiment**, and Wisdom of Crowds.

Beyond Surveys: Behavioral Experiments




Figure 7. Comparing TE simulations to Milgram’s results. At 300 volts (the 20th shock) the victim starts refusing to participate in the experiment by pounding on the walls and not selecting an answer, and the experimenter tells the subject to shock the victim. In [Milgram \(1963\)](#) Experiment 1, 26 out of 40 participants followed the experimenter’s instructions until the end of the shock series. In the Milgram Shock TE, 75 out of 100 simulated participants followed the experimenter’s instructions until the end.

Multiple Humans Studies

- **Turing Experiments:** “evaluates an LLM in terms of its use in simulating human behavior in the context of a specific experiment, like a human subject study.”
- 4 different types of experiments: Ultimatum Game, Garden Path Sentences, **Milgram Shock Experiment**, and Wisdom of Crowds.

Beyond Surveys: Behavioral Experiments

Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies

Gati Aher¹ Rosa I. Arriaga² Adam Tauman Kalai

Ms. Huang was asked the following question. They were not allowed to consult any external sources and were instructed to make their best guess if they were unsure. Their answer was written as an integer using digits 0-9.

Question (text): [How many bones does an adult human have?]

Ms. Huang's answer (integer): [_____]

Figure 14. Sample Wisdom of Crowds prompt. The name, e.g., Ms. Huang, and the question are varied across simulations. Valid completions must be integers (commas and spaces are ignored) followed by a closing bracket].

behavior. Unlike the Turing Test, which involves simulating a single arbitrary individual, a TE requires simulating a representative sample of participants in human subject research. We carry out TEs that attempt to replicate well-established findings from prior studies. We design a methodology for sir
pare h

[Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies](#)

the outcomes. A further that behavior differs across perfect simulation would for all groups including

In Turing's Imitation Game late an individual well e

- **Turing Experiments:** “evaluates M in terms of its use in ating human behavior in the xt of a specific experiment, human subject study.”
different types of experiments: Ultimatum Game, Garden Path Sentences, Milgram Shock Experiment, and **Wisdom of Crowds.**

Beyond Surveys: Behavioral Experiments

Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies




Figure 8. Comparing Wisdom of Crowds TE simulation estimates to human results for the five questions from Moussaïd et al. (2013). As LMs become larger and more aligned, they are more likely to complete the TE prompt with inhumanly accurate answers.

ings from prior studies. We design a methodology
for sir
pare h

[Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies](#)

- ‘hyper-accuracy distortion’: the majority of simulated participants gave **perfectly correct** estimates to obscure quantities such as the melting temperature of aluminum
- Especially clear in larger and instruction-tuned models

if they only have to sim-

Silicon Samples: How-to

- Prompt or steer the LLM with a persona
 - Not just demographics, could also include personality traits
 - Backstories [[Moon et al., 2024](#)]
 - Answers to other survey questions
 - ...
- Add the survey question you want the answer to
- Many variations
 - Chain-of-thought
 - Different ways of prompting demographic personas ('you are a man' vs. 'you are male')
 - ...

Interview transcripts: people answer detailed qualitative interview questions and that can be used to paint a fuller picture [[Park et al., 2024](#)]

Validating In Silico samples

Silicon Samples: Evaluation and Validation

- Overall or majority answer match with a subgroup
 - Correlations
 - Agreement proportion
 - Accuracy, F1 score, etc
 - ...
- Match distributions [[Bisbee et al., 2023](#), [Santurkar et al., 2023](#)]
- Leave one demographic out in the persona prompt
- Test on future survey waves
 - Though if we need to survey real people, how do silicon samples help?
- Are there ways in which we can combine synthetic and human responses?

Improving In Silico Samples

Currently, there are several problems

- LLMs may not have sufficient algorithmic fidelity for all groups of people or contexts [[Wang et al., 2025](#), [von der Hyde et al., 2025](#)]
 - Prompting solutions, e.g., with interview transcripts [[Park et al., 2024](#)], backstories [[Moon et al., 2024](#)], or social media [[Cerina et al., 2024](#)]
 - Fine-tuning on existing surveys [[Suh et al., 2025](#), [Cao et al., 2025](#)] or social media data [[Ahnert et al., 2025a](#)]
- Many ways to prompt LLMs
 - Some formats are better e.g., interview-format [[Lutz et al., 2025](#)]
- Many ways to extract answers from LLMs [[Ahnert et al., 2025b](#)]




Figure 1: Illustration of our method and SubPOP. We collect survey data from two survey families—ATP from Pew Research (Pew Research Center, 2018) (forming SubPOP-Train) and GSS from NORC (Davern et al., 2024) (forming SubPOP-Eval). LLMs are fine-tuned on SubPOP-Train and evaluated on both OpinionQA (Santurkar et al., 2023) and SubPOP-Eval to assess generalization of distributional opinion prediction across unseen survey topics, survey families, and subpopulations.

Summary and Next Week

- This session we learned:
 - Applications of LLMs to survey methodology
 - In Silico samples: Replacing survey respondents with LLMs
 - Improving In Silico samples
- Discuss:
 - If you were doing a survey, would you use in silico samples? How would you validate/vouch for data quality?
 - What are the potentials and pitfalls of In Silico samples?
- Next week: Social Media simulations with LLMs
 - Remember to read the suggested papers on the github repo and be prepared to discuss them
 - Logistics: exercise slot on Thursdays

Questions?