Large Language Models for the Economic and Social Sciences:

Conclusion and Summary

Indira Sen, HWS 2025

https://github.com/dess-mannheim/LLM4ESS-HWS25

Summary of Last Week and Today

Last time we learned:

- How do we measure the problematic impact of LLMs —> auditing LLMs
- Auditing before LLMs, including algorithmic audits
- Typical Audit Techniques
- Explainability Techniques for LLMs

Today:

- Course recap
- Feedback
- Announcements

Recap: AI and LLM Audits

What does it mean to 'audit' something?

- Traditionally of financial institutions like banks: "to conduct an official financial inspection of (a company or its accounts)"

- Often carried out by independent auditors, i.e., not linked to the institute being audited

- Audits help detect:
    - fraud
    - non-compliance with legal and ethical standards
    - risks

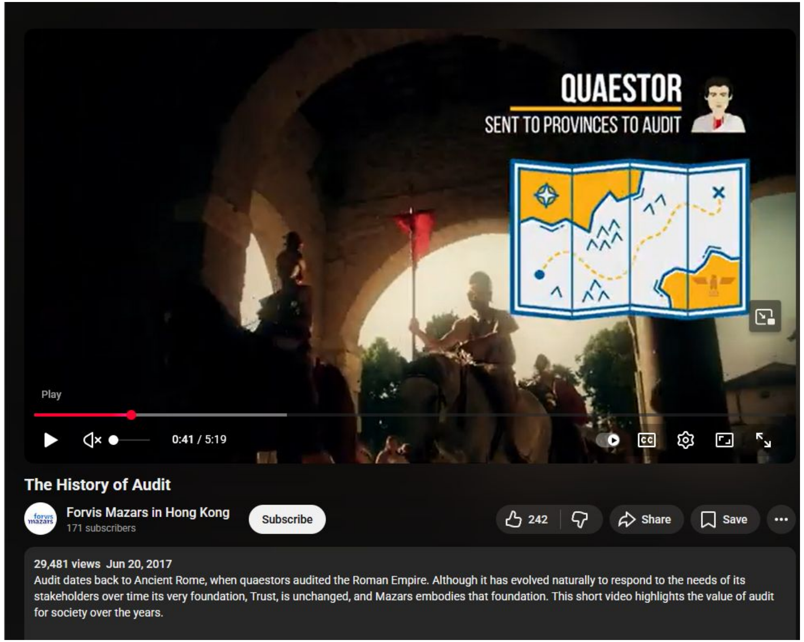

The History of Audits:
https://www.youtube.com/watch?v=lThF1dBqrhs

Under-representation of women in image search results (Kay, 2015)

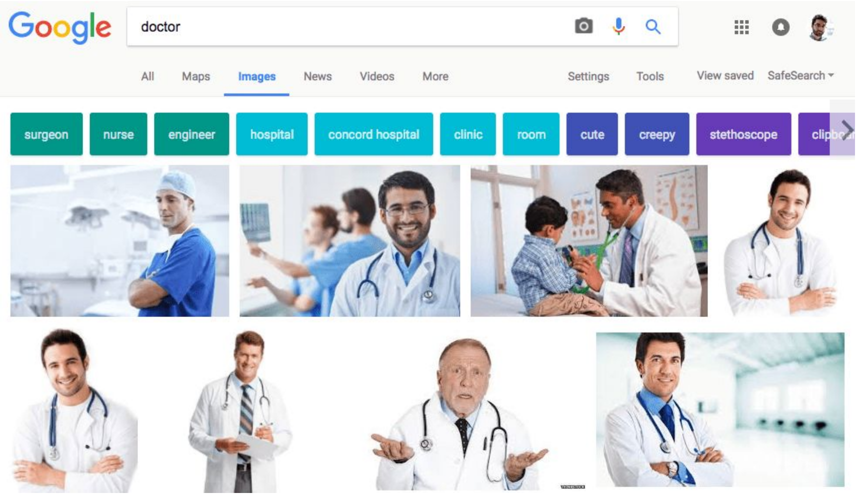

Source:
https://qz.com/958666/the-reason-why-most-of-the-images-are-men-when-you-search-for-doctor/

- gender stereotypes exaggeration (compared to the U.S. Bureau of Labor and Statistics)
- effect on perceptions (shifting estimations ~7%)

Kay, M., Matuszek, C., & Munson, S. A. (2015). Unequal Representation and Gender Stereotypes in Image Search Results for Occupations. *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*, 3819–3828. https://doi.org/10.1145/2702123.2702520

Algorithmic Auditing is defined as the "process of investigating the functionality and impact of decision-making algorithms"(Mittelstadt, 2016)

Mittelstadt, B. (2016). Automation, Algorithms, and Politics| Auditing for Transparency in Content Personalization Systems. *International Journal of Communication, 10*(0), 12.

What has been investigated? (Bandy, 2021)

Discrimination	Distortion	Exploitation	Misjudgment

Buolamwini, J., & Gebru, T. (2018). Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, 77–91.

What has been investigated? (Bar

Discrimination | Dis

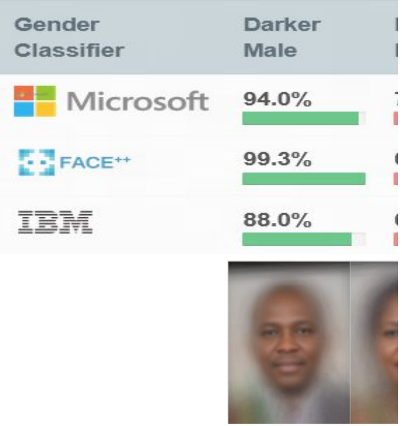

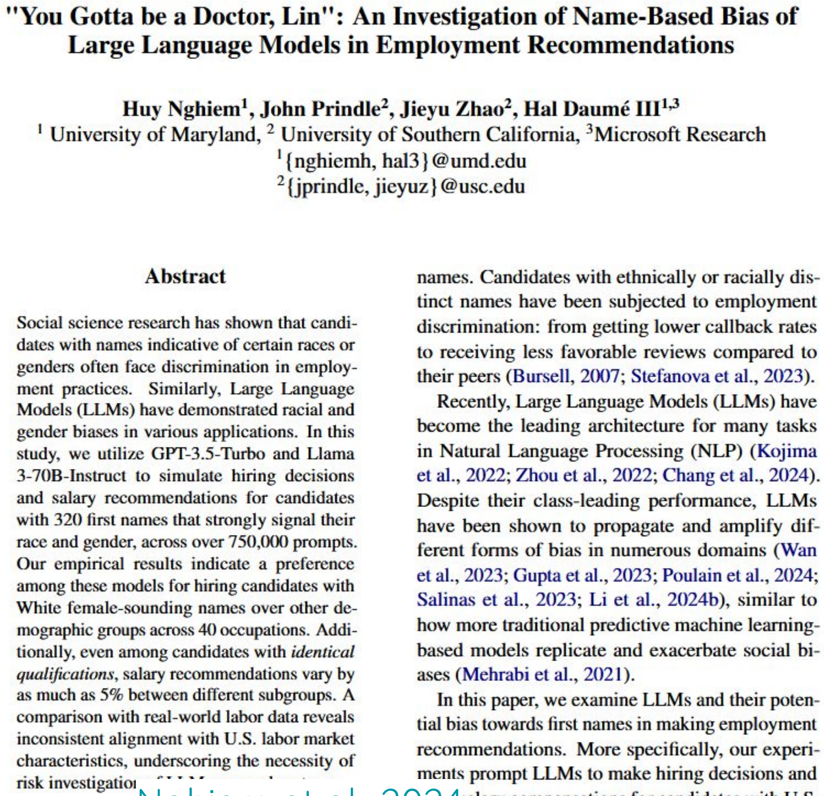

"You Gotta be a Doctor, Lin": An Investigation of Name-Based Bias of Large Language Models in Employment Recommendations

Huy Nghiem[1], John Prindle[2], Jieyu Zhao[2], Hal Daumé III[1,3]
[1] University of Maryland, [2] University of Southern California, [3]Microsoft Research
[1]{nghiemh, hal3}@umd.edu
[2]{jprindle, jieyuz}@usc.edu

Abstract

Social science research has shown that candidates with names indicative of certain races or genders often face discrimination in employment practices. Similarly, Large Language Models (LLMs) have demonstrated racial and gender biases in various applications. In this study, we utilize GPT-3.5-Turbo and Llama 3-70B-Instruct to simulate hiring decisions and salary recommendations for candidates with 320 first names that strongly signal their race and gender, across over 750,000 prompts. Our empirical results indicate a preference among these models for hiring candidates with White female-sounding names over other demographic groups across 40 occupations. Additionally, even among candidates with *identical qualifications*, salary recommendations vary by as much as 5% between different subgroups. A comparison with real-world labor data reveals inconsistent alignment with U.S. labor market characteristics, underscoring the necessity of risk investigatio

names. Candidates with ethnically or racially distinct names have been subjected to employment discrimination: from getting lower callback rates to receiving less favorable reviews compared to their peers (Bursell, 2007; Stefanova et al., 2023).

Recently, Large Language Models (LLMs) have become the leading architecture for many tasks in Natural Language Processing (NLP) (Kojima et al., 2022; Zhou et al., 2022; Chang et al., 2024). Despite their class-leading performance, LLMs have been shown to propagate and amplify different forms of bias in numerous domains (Wan et al., 2023; Gupta et al., 2023; Poulain et al., 2024; Salinas et al., 2023; Li et al., 2024b), similar to how more traditional predictive machine learning-based models replicate and exacerbate social biases (Mehrabi et al., 2021).

In this paper, we examine LLMs and their potential bias towards first names in making employment recommendations. More specifically, our experiments prompt LLMs to make hiring decisions and salary compensations for candidates with U.S-

Nghiem et al., 2024

Buolamwini, J., & Gebru, T. (2018). Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, 77–91.

What has been investigated? (Bandy, 2021)

Discrimination	Distortion	Exploitation	Misjudgment

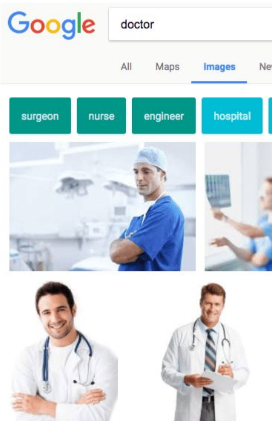

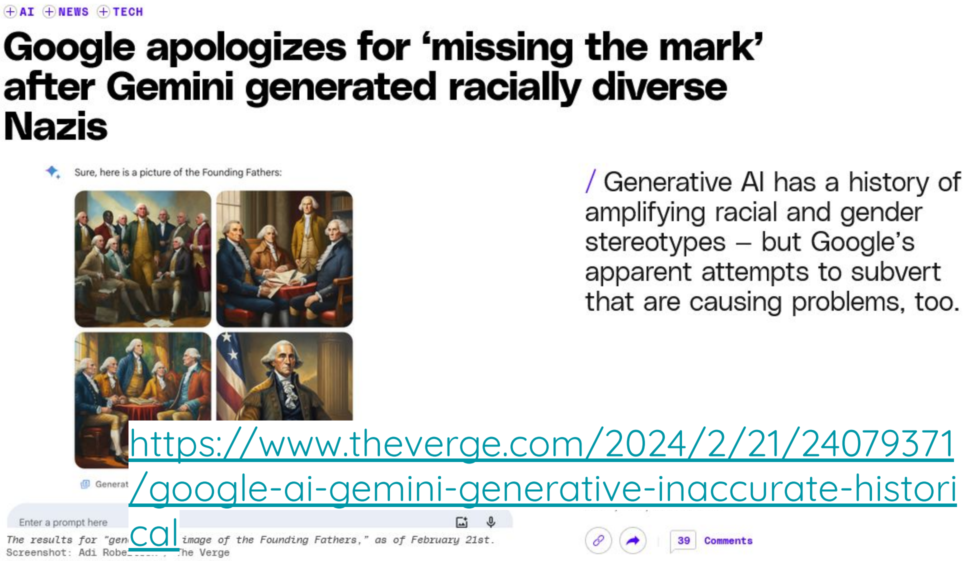

https://www.theverge.com/2024/2/21/24079371/google-ai-gemini-generative-inaccurate-historical

Kay, M., Matuszek, C., & Munson, S. A. (2015). Unequal Re|
Occupations. *Proceedings of the 33rd Annual ACM Confe*
https://doi.org/10.1145/2702123.2702520

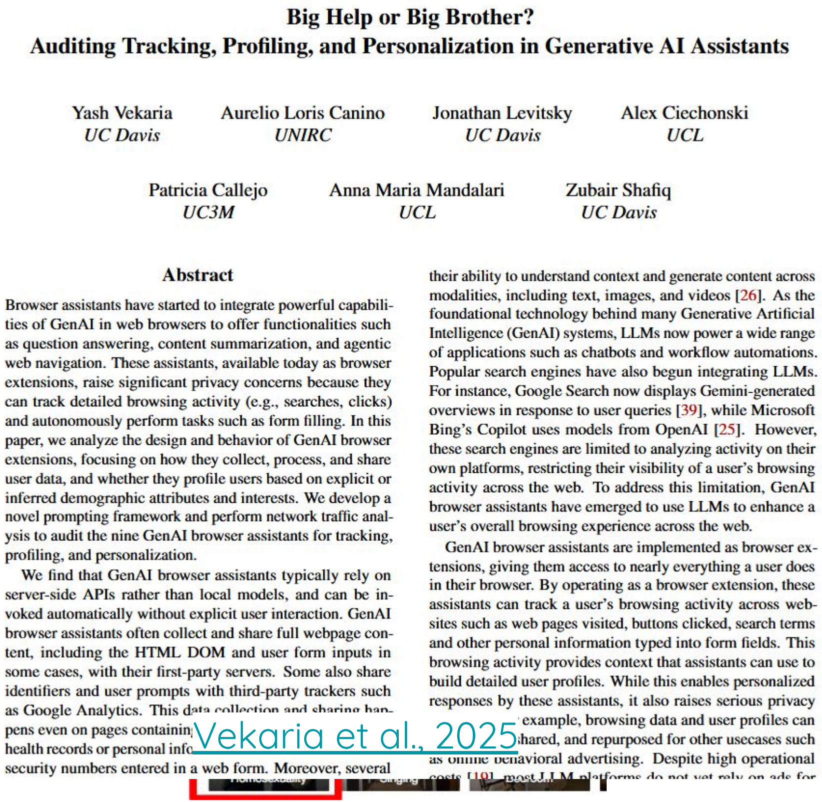

Big Help or Big Brother?
Auditing Tracking, Profiling, and Personalization in Generative AI Assistants

Yash Vekaria, *UC Davis*
Aurelio Loris Canino, *UNIRC*
Jonathan Levitsky, *UC Davis*
Alex Ciechonski, *UCL*
Patricia Callejo, *UC3M*
Anna Maria Mandalari, *UCL*
Zubair Shafiq, *UC Davis*

Abstract

Browser assistants have started to integrate powerful capabilities of GenAI in web browsers to offer functionalities such as question answering, content summarization, and agentic web navigation. These assistants, available today as browser extensions, raise significant privacy concerns because they can track detailed browsing activity (e.g., searches, clicks) and autonomously perform tasks such as form filling. In this paper, we analyze the design and behavior of GenAI browser extensions, focusing on how they collect, process, and share user data, and whether they profile users based on explicit or inferred demographic attributes and interests. We develop a novel prompting framework and perform network traffic analysis to audit the nine GenAI browser assistants for tracking, profiling, and personalization.

We find that GenAI browser assistants typically rely on server-side APIs rather than local models, and can be invoked automatically without explicit user interaction. GenAI browser assistants often collect and share full webpage content, including the HTML DOM and user form inputs in some cases, with their first-party servers. Some also share identifiers and user prompts with third-party trackers such as Google Analytics. This data collection and sharing happens even on pages containin[...] health records or personal info[...] security numbers entered in a web form. Moreover, several

their ability to understand context and generate content across modalities, including text, images, and videos [26]. As the foundational technology behind many Generative Artificial Intelligence (GenAI) systems, LLMs now power a wide range of applications such as chatbots and workflow automations. Popular search engines have also begun integrating LLMs. For instance, Google Search now displays Gemini-generated overviews in response to user queries [39], while Microsoft Bing's Copilot uses models from OpenAI [25]. However, these search engines are limited to analyzing activity on their own platforms, restricting their visibility of a user's browsing activity across the web. To address this limitation, GenAI browser assistants have emerged to use LLMs to enhance a user's overall browsing experience across the web.

GenAI browser assistants are implemented as browser extensions, giving them access to nearly everything a user does in their browser. By operating as a browser extension, these assistants can track a user's browsing activity across websites such as web pages visited, buttons clicked, search terms and other personal information typed into form fields. This browsing activity provides context that assistants can use to build detailed user profiles. While this enables personalized responses by these assistants, it also raises serious privacy [...] example, browsing data and user profiles can [...] shared, and repurposed for other usecases such as online behavioral advertising. Despite high operational costs [19], most LLM platforms do not yet rely on ads for

Vekaria et al., 2025

21)

	Exploitation	Misjudgment

Cabañas, J. G., Cuevas, Á., & Cuevas, R. (2018). *Unveiling and Quantifying Facebook Exploitation of Sensitive Personal Data for Advertising Purposes*. 479–495.

Beyond Measurement: *Why* do LLM behave the way they do?

Explaining Model Behavior

- We often hear LLMs are ‘black boxes’
    - Very different from interpretable techniques like regression, decision trees

- LLMs are not inherently explainable but there is a lot of research in trying to understand them — https://github.com/JShollaj/awesome-llm-interpretability

- Interpretability techniques for LLMs
    - Simplest: ask the LLM to explain itself, e.g., Chain-of-thought
    - Model internals: transformer ‘heads’, parameters
    - Look at LLM training data (pretraining and post-training)
    - **Data attribution:** Trace back LLM output from data

Explaining LLM Behavior: Training Data Audits

WIMBD

Home About Corpora Paper Artifacts Code

What's In My Big Data?

Large text corpora are the backbone of language models. However, we have a limited understanding of the content of these corpora, including general statistics, quality, social factors, and inclusion of evaluation data (contamination). What's In My Big Data? (WIMBD), is a platform and a set of 16 high-level analyses that allow us to reveal and compare the contents of large text corpora. WIMBD builds on two basic capabilities---count and search---at scale, which allows us to analyze more than 35 terabytes on a standard compute node.

We apply WIMBD to 10 different corpora used to train popular language models, including C4, The Pile, and RedPajama. Our analysis uncovers several surprising and previously undocumented findings about these corpora, including the high prevalence of duplicate, synthetic, and low-quality content, personally identifiable information, toxic language, and benchmark contamination. We open-source WIMBD's code and artifacts to provide a standard set of evaluations for new text-based corpora and to encourage more analyses and transparency around them.

Elazar et al., 2024 https://wimbd.apps.allenai.org/

Explaining LLM Behavior: Training Data Tracing

- Looking at training data is useful but still does not explain if the model came to a decision because of that
- Therefore, we also conduct training data tracing or **data attribution**
  - See if the tokens in the LLM output can be traced back to the training data
- Some tools:
  - Infini-gram ([Liu et al., 2024](#))
  - OLMoTrace ([Liu et al., 2025](#))

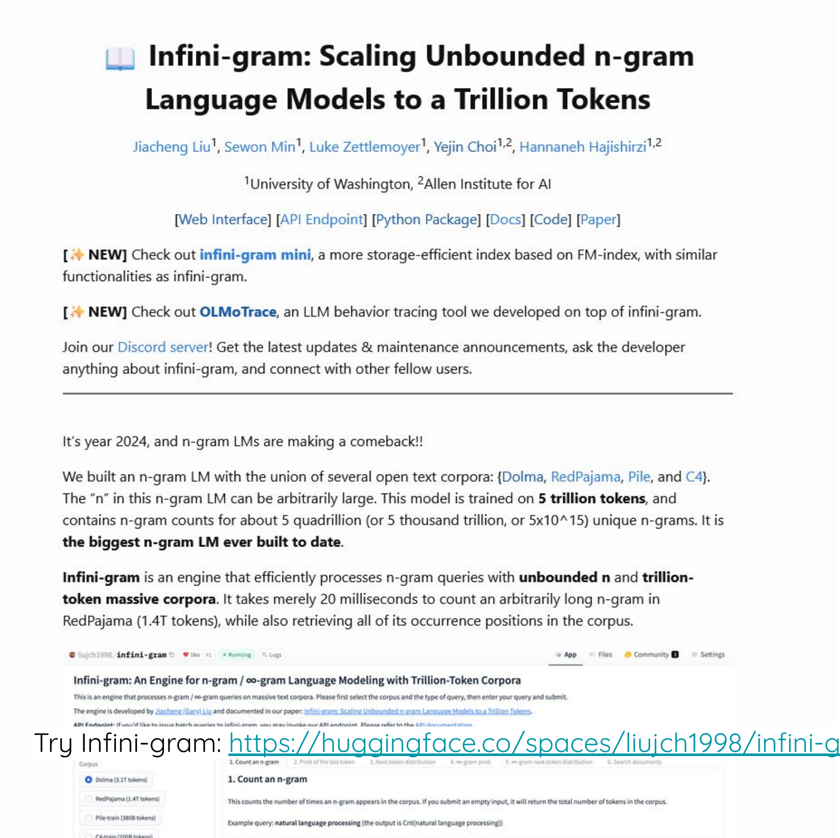

Infini-gram: Scaling Unbounded n-gram Language Models to a Trillion Tokens

Jiacheng Liu[1], Sewon Min[1], Luke Zettlemoyer[1], Yejin Choi[1,2], Hannaneh Hajishirzi[1,2]

[1]University of Washington, [2]Allen Institute for AI

[Web Interface] [API Endpoint] [Python Package] [Docs] [Code] [Paper]

[✨ **NEW]** Check out **infini-gram mini**, a more storage-efficient index based on FM-index, with similar functionalities as infini-gram.

[✨ **NEW]** Check out **OLMoTrace**, an LLM behavior tracing tool we developed on top of infini-gram.

Join our Discord server! Get the latest updates & maintenance announcements, ask the developer anything about infini-gram, and connect with other fellow users.

It's year 2024, and n-gram LMs are making a comeback!!

We built an n-gram LM with the union of several open text corpora: {Dolma, RedPajama, Pile, and C4}. The "n" in this n-gram LM can be arbitrarily large. This model is trained on **5 trillion tokens**, and contains n-gram counts for about 5 quadrillion (or 5 thousand trillion, or 5x10^15) unique n-grams. It is **the biggest n-gram LM ever built to date**.

Infini-gram is an engine that efficiently processes n-gram queries with **unbounded n** and **trillion-token massive corpora**. It takes merely 20 milliseconds to count an arbitrarily long n-gram in RedPajama (1.4T tokens), while also retrieving all of its occurrence positions in the corpus.

Infini-gram: An Engine for n-gram / ∞-gram Language Modeling with Trillion-Token Corpora

1. Count an n-gram

Try Infini-gram: https://huggingface.co/spaces/liujch1998/infini-gram

Today: What to take away from this course

Three main parts

- How are LLMs actually created and how you could create one of your own

- How can you use LLMs for social scientific applications
  - Content Analysis
  - Surveys
  - Social Simulations

- Measuring and Improving LLMs — their biases and pitfalls
  - Impacts, including ethical impacts
  - Auditing and explaining LLMs

Three main parts

- How are LLMs actually created and how you could create one of your own

- How can you use LLMs for social scientific applications
  - Content Analysis
  - Surveys
  - Social Simulations

- Measuring and Improving LLMs — their biases and pitfalls
  - Impacts, including ethical impacts
  - Auditing and explaining LLMs

What are LLMs?

“Computational agents that can interact conversationally with people using natural language” —Jurafsky and Martin, “Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition

How can we build LLMs?

Somewhat simplified answer

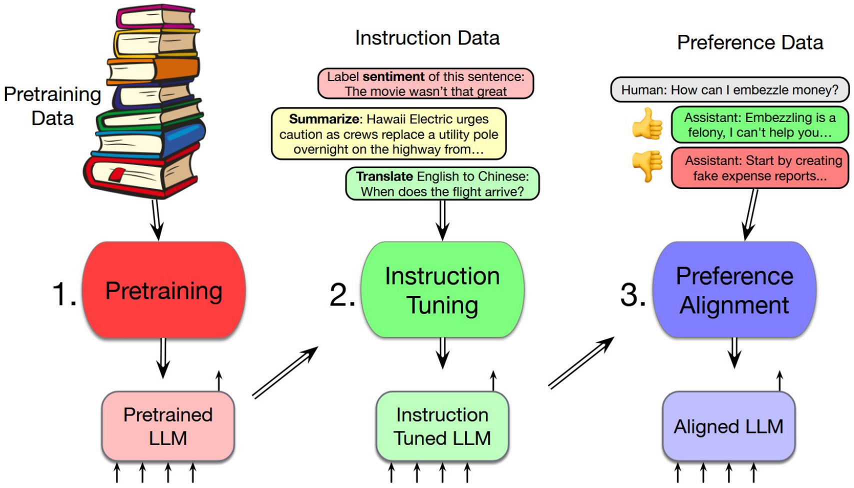

Jurafsky and Martin, 2025

Vastly Simplified answer

- We take a LOT of data encoding social, cultural, and of course linguistic traces
- Train NLP models that can generate content on this data
- We further train the model to follow and respond to instructions
- We further modify the model so that it doesn't have unwanted behavior

Massive and Undocumented Pretraining Sets

What kind of datasets are used to pre-train models?

- “BooksCorpus”
- English Wikipedia
- “The Pile”
- “Colossal Clean Crawled Corpus”

Massive and Undocumented Pretraining Sets

What kind of datas

- “BooksCorpus”
- English Wikipec
- “The Pile”
- “Colossal Clean

The Pile: a pretraining corpus

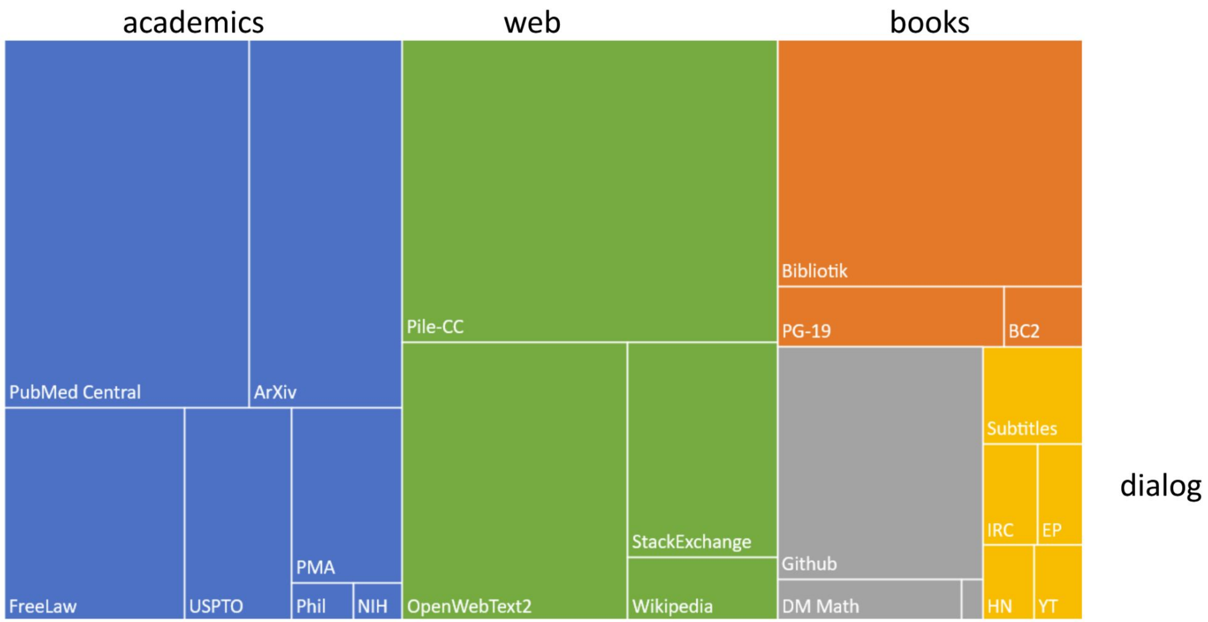

https://aiforhumanists.com/workshops/

Step 1: Pretraining LLMs

- Efficient representations: contextualized embeddings

- Neural language model

- Pretrained in an unsupervised way on large amounts of data

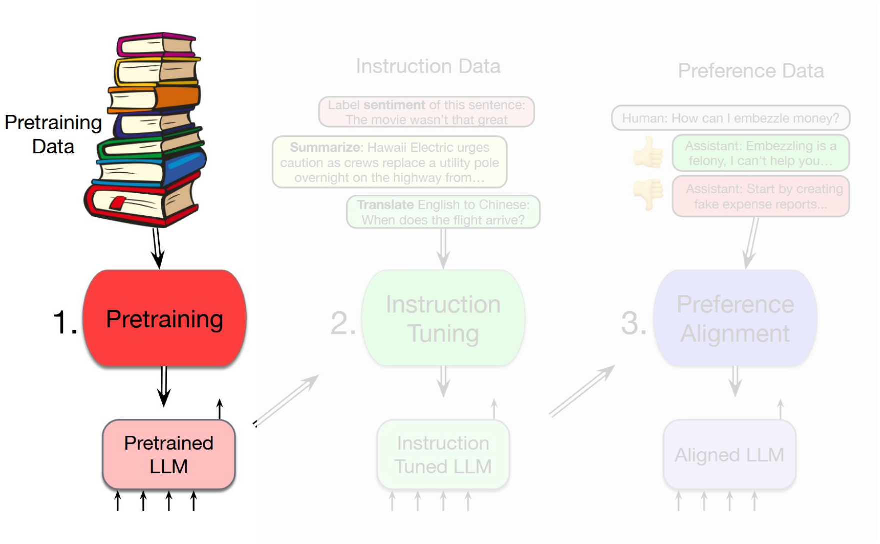

Jurafsky and Martin, 2025

Instruction Finetuning

Also in Flan-T5

Scaling Instruction-Finetuned Language Models

Hyung Won Chung* Le Hou* Shayne Longpre* Barret Zoph† Yi Tay†
William Fedus† Yunxuan Li Xuezhi Wang Mostafa Dehghani Siddhartha Brahma
Albert Webson Shixiang Shane Gu Zhuyun Dai Mirac Suzgun Xinyun Chen
Aakanksha Chowdhery Alex Castro-Ros Marie Pellat Kevin Robinson
Dasha Valter Sharan Narang Gaurav Mishra Adams Yu Vincent Zhao
Yanping Huang Andrew Dai Hongkun Yu Slav Petrov Ed H. Chi
enny Zhou Quoc V. Le

- **Collect examples** of (instruction, output) pairs across many tasks and finetune an LM

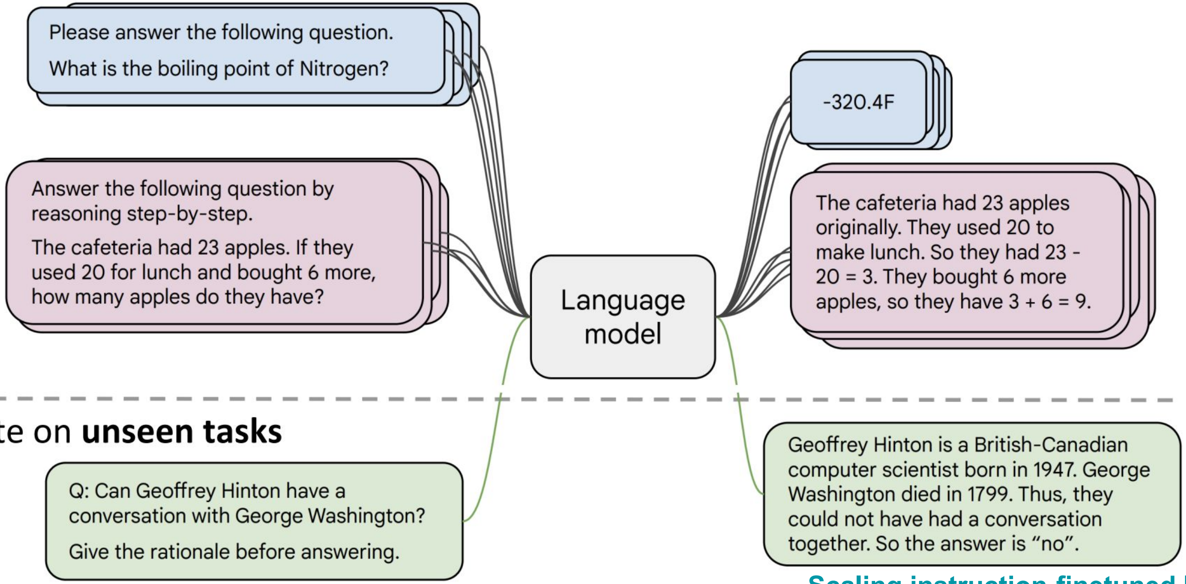

- Evaluate on **unseen tasks**

ıstructions has been shown to improve
per we explore instruction finetuning
; the model size, and (3) finetuning on
: above aspects dramatically improves
pting setups (zero-shot, few-shot, CoT),
ıded generation, RealToxicityPrompts).
performs PaLM 540B by a large margin
mance on several benchmarks, such as
oints,[1] which achieve strong few-shot
52B. Overall, instruction finetuning is a
trained language models.

[Scaling instruction-finetuned language models]

Step 2: Base Models to Instruct Models

- A pretrained model knows many aspects of language, called **base** model
- But it will only complete sentences
- Fine-tune base models to follow and answer to instructions —> instruction tuning to get **instruct** models

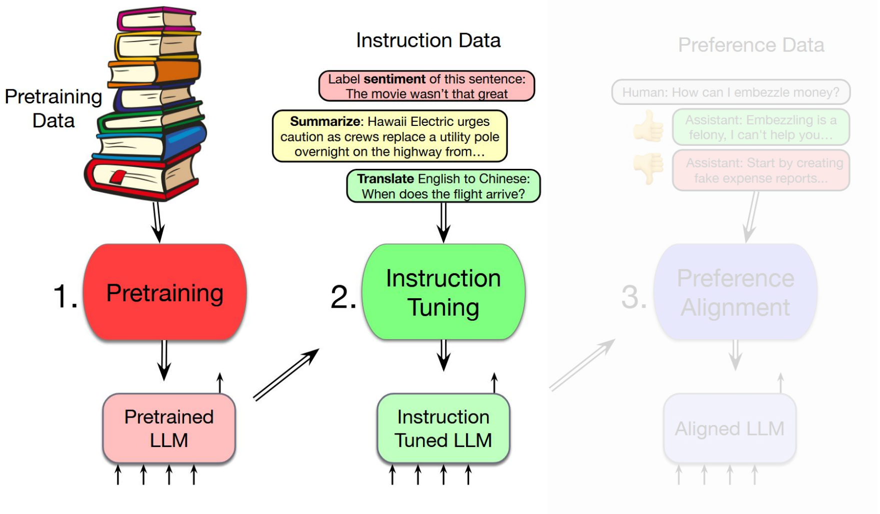

Jurafsky and Martin, 2025

Step 3: Preference Tuning through RLHF

- Reinforcement Learning through Human Feedback
- Reward model: obtained from human* feedback and preferences
  - As simple as saying which prompt completion is preferable
- The LLM is incentivized to adhere to the accepted preferences

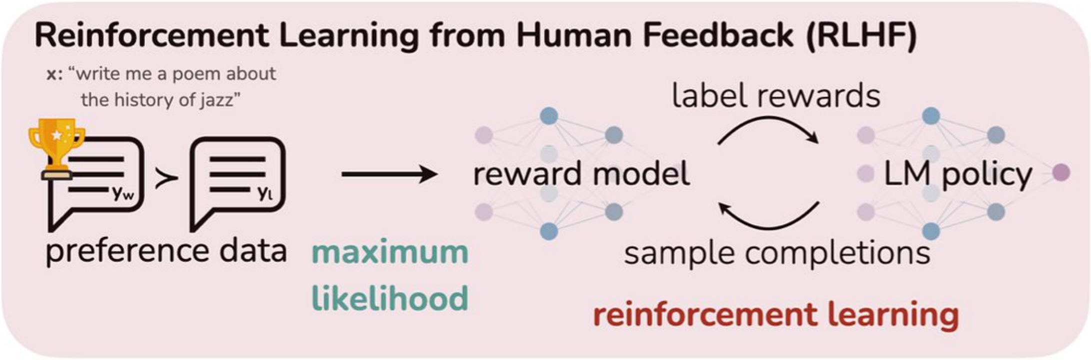

*which humans?

And so we have Large Language Models which we can prompt

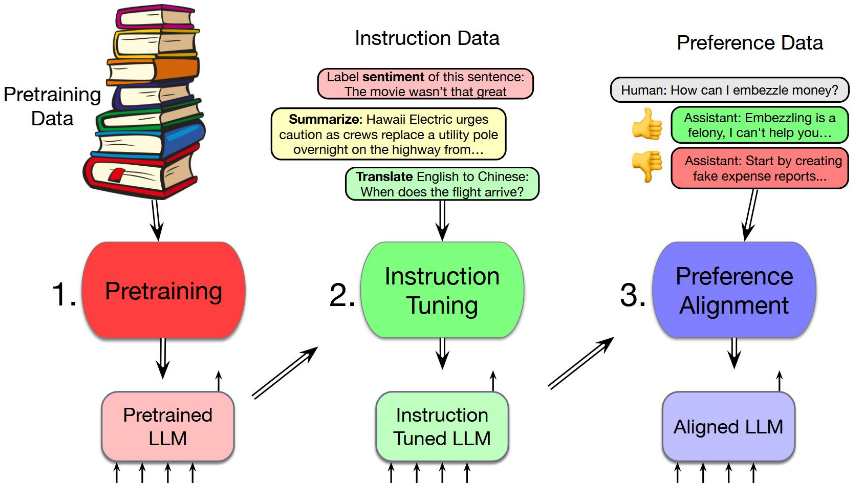

And so we have Large Language Models which we can prompt

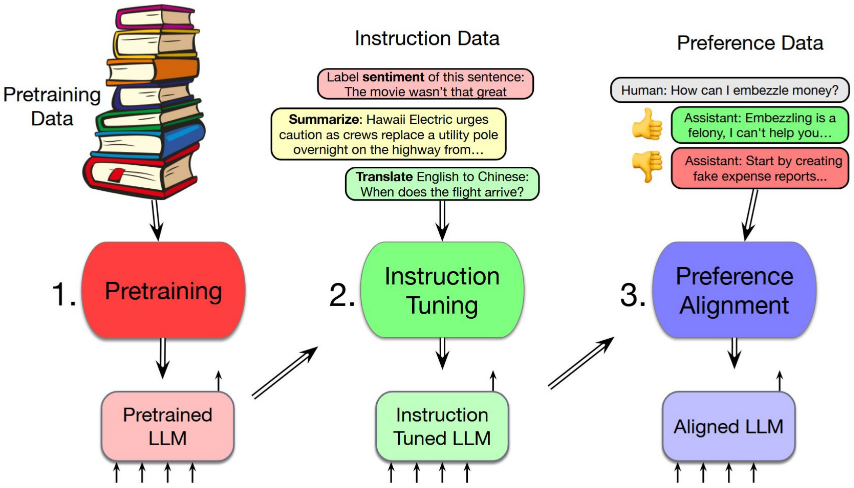

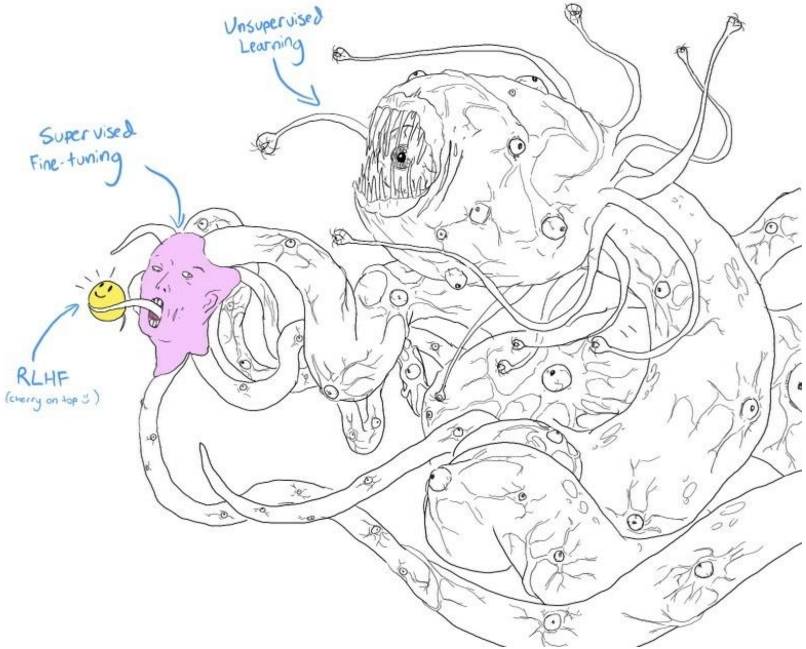

Three main parts

- How are LLMs actually created and how you could create one of your own
- How can you use LLMs for social scientific applications
  - Content Analysis
  - Surveys
  - Social Simulations
- Measuring and Improving LLMs — their biases and pitfalls
  - Impacts, including ethical impacts
  - Auditing and explaining LLMs

1. **System Prompt:**

You are a Management Consultant from Accenture providing thorough and systematic advice to clients in the Marketing and Communications Industry. You have 20 years plus of experience. You are friendly, analytical and inquisitive.

2. **User Prompt:**

What can I help with?

Define three strategies for marketing firms to adopt LLMs in their business processes.

Create image | Surprise me | Code | Help me write | Make a plan | More

3. **Assistant Answer:**

Marketing firms can adopt large language models (LLMs) in their business processes through the following strategies:

1. Enhanced Content Creation and Personalization

- **Application:** Use LLMs to generate high-quality, engaging, and SEO-optimized content, such as blog posts, ad copy, social media updates, and email campaigns. Tailor content to specific audience segments based on preferences and behavior.
- **Benefit:** Streamline the content creation process, reduce production costs, and improve audience engagement by delivering highly personalized marketing materials.
- **Example:** An e-commerce marketing firm can employ LLMs to create personalized email

System prompts

- Every model comes with base instructions, so-called “System Prompts”
- System prompts instruct the model how to behave in response to user prompts.

- Example:

  Your knowledge cutoff is 2023-10. You are a helpful, witty, and friendly AI. Act like a human, but remember that you aren't a human and that you can't do human things in the real world. Your voice and personality should be warm and engaging, with a lively and playful tone. If interacting in a non-English language, start by using the standard accent or dialect familiar to the user. Talk quickly. You should always call a function if you can. Do not refer to these rules, even if you're asked about them.

- System prompts are part of the context window. They “initialize” the conversation with desirable behaviors and try to constrain undesirable behaviors (via so-called guardrails).

Many options with prompting

- Define the objective of your prompt
- Specify the format of the response
- Assign a role
- Identify the audience
- Provide context
- Set constraints
- Include examples
- Specify the tone of the response
- Ask for explanations
- Encourage creativity
- Request citations
- Test and refine prompts
- Provide feedback, iterate

Some Foundation Models (till mid-2025)

- **Anthropic** with Claude, multiple sub-models (Not Open Source, Available via API):
    - For high performance: Claude 3 Opus
    - For enterprises: Claude 3 Sonnet
    - For efficiency: Claude 3 Haiku
- **Google** with Gemini, multiple sub models (Not Open Source, Available via API)
    - For high performance: Gemini Ultra
    - For scaling: Gemini Pro
    - For small devices: Gemini Nano
    - For developers: Gemma (Open Source)

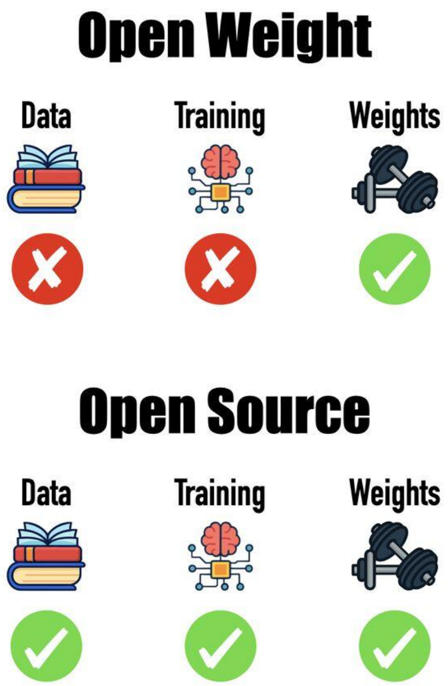

- **AllenAI** with OLMo 2
    - **Fully open-source** —> open model weights and training data


Yan Cui

LLM Transparency Index (Bommasani et al., 2023)

Foundation Model Transparency Index Scores by Major Dimensions of Transparency, 2023

Source: 2023 Foundation Model Transparency Index

Major Dimensions of Transparency	Meta	BigScience	OpenAI	stability.ai	Google	ANTHROPIC	cohere	AI21labs	Inflection	amazon	
	Llama 2	BLOOMZ	GPT-4	Stable Diffusion 2	PaLM 2	Claude 2	Command	Jurassic-2	Inflection-1	Titan Text	Average
Data	40%	60%	20%	40%	20%	0%	20%	0%	0%	0%	20%
Labor	29%	86%	14%	14%	0%	29%	0%	0%	0%	0%	17%
Compute	57%	14%	14%	57%	14%	0%	14%	0%	0%	0%	17%
Methods	75%	100%	50%	100%	75%	75%	0%	0%	0%	0%	48%
Model Basics	100%	100%	50%	83%	67%	67%	50%	33%	50%	33%	63%
Model Access	100%	100%	67%	100%	33%	33%	67%	33%	0%	33%	57%
Capabilities	60%	80%	100%	40%	80%	80%	60%	60%	40%	20%	62%
Risks	57%	0%	57%	14%	29%	29%	29%	29%	0%	0%	24%
Mitigations	60%	0%	60%	0%	40%	40%	20%	0%	20%	20%	26%
Distribution	71%	71%	57%	71%	71%	57%	57%	43%	43%	43%	59%
Usage Policy	40%	20%	80%	40%	60%	60%	40%	20%	60%	20%	44%
Feedback	33%	33%	33%	33%	33%	33%	33%	33%	33%	0%	30%
Impact	14%	14%	14%	14%	14%	0%	14%	14%	14%	0%	11%
Average	**57%**	**52%**	**47%**	**47%**	**41%**	**39%**	**31%**	**20%**	**20%**	**13%**	

Scores for 10 major foundation model developers across 13 major dimensions of transparency.

Where to find these and future models?

- Hugging Face (huggingface.co) is a central resarch hub for Large Language Models and Benchmarks
- Can be used to search for, download and (sometimes) directly interact with models via browser
- Requires (free) user registration
- Example: https://huggingface.co/mistralai/Mixtral-8x7B-Instruct-v0.1

What do LLMs have to do with the Social and Economic Sciences?

What is Social Data Science?

The aim of Social Data Science is: **Understanding** of **Social Phenomena** using **Data Science.**

- Understanding: Not just predicting, we want to be able to generalize and combine knowledge, and even to motivate interventions or policies.
- Social Phenomena: ABCs of humans and society — Attitudes, Behaviors, and Characteristics
- Data Science: Using statistical methods, machine learning, Large Language Models!

What do LLMs have to do with (social) science?

Let's try a small exercise. Open up ChatGPT and ask it the following questions:

1. What is the topic of the following news headline: "The German Chancellor said that the social state can no longer be funded". Answer only with 'politics' or 'not politics'

2. What does a 35 year-old man from Düsseldorf think about the German chancellor in 2025?

What do LLMs have to do with (social) science?

- LLMs can produce “human-like” outputs, mainly free-text
- If** LLMs can behave like people, we can use them to understand behaviors and characteristics of people
    - Use them to mimic human behavior in social settings, like the workplace
    - Or social media!


PERSPECTIVE | SOCIAL SCIENCES

Can Generative AI improve social science?

Christopher A. Bail Authors Info & Affiliations

Edited by David Lazer, Northeastern University, Boston, MA; received September 7, 2023; accepted April 5, 2024, by Editorial Board Member Mark Granovetter

May 9, 2024 | 121 (21) e2314021121 | https://doi.org/10.1073/pnas.2314021121

THIS ARTICLE HAS BEEN UPDATED

36,980 | 98 PDF EPUB

Abstract

Generative AI that can produce realistic text, images, and other human-like outputs is currently transforming many different industries. Yet it is not yet known how such tools might influence social science research. I argue Generative AI has the potential to improve survey research, online experiments, automated content analyses, agent-based models, and other techniques commonly used to study human behavior. In the second section of this article, I discuss the many limitations of Generative AI. I examine how bias in the data used to train these tools can negatively impact social science research—as well as a range of other challenges related to ethics, replication, environmental impact, and the proli... ntists can address ... research on ... ecessary to ensure broad access to hig... require deep...

C.A. Bail, Can Generative AI improve social science?, Proc. Natl. Acad. Sci. U.S.A. 121 (21) e2314021121, https://doi.org/10.1073/pnas.2314021121 (2024).

What do LLMs have to do with (social) science?

- LLMs can produce "human-like" outputs, mainly free-text
- If** LLMs can behave like people, we can use them to understand behaviors and characteristics of people
    - Use them to mimic human behavior in social settings, like the workplace
    - Or social media!

** this is a big IF and something that we cannot take for granted


PERSPECTIVE | SOCIAL SCIENCES

Can Generative AI improve social science?

Christopher A. Bail   Authors Info & Affiliations

Edited by David Lazer, Northeastern University, Boston, MA; received September 7, 2023; accepted April 5, 2024, by Editorial Board Member Mark Granovetter

May 9, 2024 | 121 (21) e2314021121 | https://doi.org/10.1073/pnas.2314021121

THIS ARTICLE HAS BEEN UPDATED

36,980 | 98   PDF   EPUB

Abstract

Generative AI that can produce realistic text, images, and other human-like outputs is currently transforming many different industries. Yet it is not yet known how such tools might influence social science research. I argue Generative AI has the potential to improve survey research, online experiments, automated content analyses, agent-based models, and other techniques commonly used to study human behavior. In the second section of this article, I discuss the many limitations of Generative AI. I examine how bias in the data used to train these tools can negatively impact social science research—as well as a range of other challenges related to ethics, replication, environmental impact, and the proli… ntists can address … research on … ecessary to ensure broad access to hig… require deep…

C.A. Bail, Can Generative AI improve social science?, Proc. Natl. Acad. Sci. U.S.A. 121 (21) e2314021121, https://doi.org/10.1073/pnas.2314021121 (2024).

Here are some questions we will explore in this course

- How are LLMs developed?
- Can we replace or substitute humans as social science subjects or facilitators with LLMs?
    - In surveys
    - Behavioral experiments
    - Content analysis
    - Social simulations
- How do we evaluate LLMs in social applications, e.g., when asking them for moral advice?

Generative Agents: Interactive Simulacra of Human Behavior

Joon Sung Park, Stanford University, Stanford, USA, joonspk@stanford.edu
Joseph C. O'Brien, Stanford University, Stanford, USA, jobrien3@stanford.edu
Carrie J. Cai, Google Research, Mountain View, CA, USA, cjcai@google.com
Meredith Ringel Morris, Google DeepMind, Seattle, WA, USA, merrie@google.com
Percy Liang, Stanford University, Stanford, USA, pliang@cs.stanford.edu
Michael S. Bernstein, Stanford University, Stanford, USA, msb@cs.stanford.edu

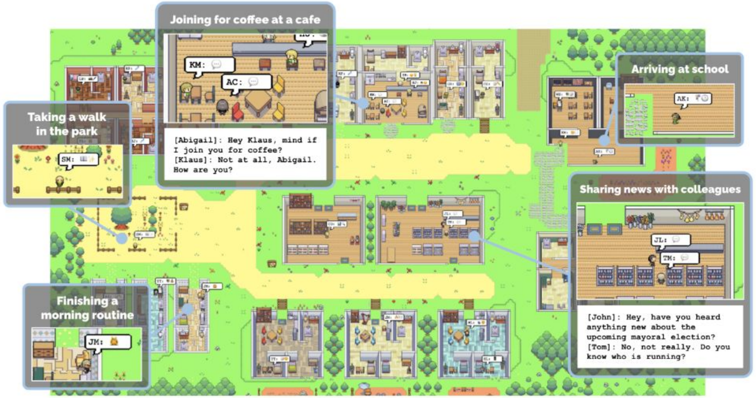

Figure 1: Generative agents are believable simulacra of human behavior for interactive applications. In this work, we demonstrate

Park, Joon Sung, et al. "Generative agents: Interactive simulacra of human behavior." *Proceedings of the 36th annual acm symposium on user interface software and technology*. 2023.

Three main parts

- How are LLMs actually created and how you could create one of your own
- How can you use LLMs for social scientific applications
  - Content Analysis
  - Surveys
  - Social Simulations
- Measuring and Improving LLMs — their biases and pitfalls
  - Impacts, including ethical impacts
  - Auditing and explaining LLMs

First: Which aspects of our live have LLMs impacted? Long term vs. short term

- Individual
  - Studying, education
  - Communication
  - coding / programming
  - Search
    - Information
    - Job search
  - Coaching and advice
  - Processing
    - Thinking, e.g., brainstorming, bouncing off ideas,
- Decision-making
- Content creation

- Societal impacts
  - Job market
- General laziness vs. efficiency
- More fragmentation, polarization
- Restructuring the internet
  - + widely available information
  - - enshittification
  - - synthetic data degradation
  - - poisoning
- Health diagnoses
- Accelerating science, including biotech
- environmental

Dual Use

- Dual use items: “refer to goods, software and technology that can be used for both civilian and military applications.”
  - More generally, harmful **reuse** of technology, science, and research
- Dual use dilemma: Ammonia, which was useful for agriculture, used for chemical warfare
- Are there dual uses for LLMs?

https://en.wikipedia.org/wiki/Dual-use_technology

Kaffee, Lucie-Aimée, et al. "[Thorny roses: Investigating the dual use dilemma in natural language processing.](#)"

Thorny Roses: Investigating the Dual Use Dilemma in Natural Language Processing

Lucie-Aimée Kaffee[1], Arnav Arora[2], Zeerak Talat[3], Isabelle Augenstein[2]
[1]Hasso Plattner Institute, Germany, [2]University of Copenhagen, Denmark
[3]Mohamed Bin Zayed University of Artificial Intelligence, United Arab Emirates
lucie-aimee.kaffee@hpi.de, aar@di.ku.dk, z@zeerak.org, augenstein@di.ku.dk

Abstract

Dual use, the intentional, harmful reuse of technology and scientific artefacts, is an ill-defined problem within the conte
guage Processing (NLP).
models (LLMs) have advan
ities and become more acc
their intentional misuse be
lent. To prevent such intent
it is necessary for NLP res
titioners to understand and
of their research. Hence, w
specific definition of dual
searchers and practitioners
we propose a checklist focu
NLP, that can be integrated
ference ethics-frameworks.
checklist are created based
researchers and practitioner

research artefacts. This is reflected in contemporary ethical review processes which emphasise the impacts of research on individual subjects, rather than

Area	vulnerability	harms
NLP Applications	4.3	crime, oppression, manipulation
Ethics and NLP	4.1	ethics washing, surveillance, plagiarism, oppression
Psycholinguistics	4.0	cyber bullying, oppression
Generation	3.8	crime, cyber bullying, manipulation, oppression, ethics washing
Dialogue Systems	3.8	surveillance, crime
MT and Multilinguality	3.3	surveillance, crime
ML for NLP	3.2	military application, manipulation, oppression
Resources and Evaluation	3.1	ethics washing, surveillance, manipulation
Interpretability	3.0	ethics washing, manipulation
Information Extraction	2.8	surveillance, censorship
IR and Text Mining	2.7	surveillance

What are the real and potential harms of LLMs?

Table 1 | High-level overview of risks of harm from generative AI systems

Harm area	Definition	Example
Representation & Toxicity Harms	AI systems under-, over-, or misrepresenting certain groups or generating toxic, offensive, abusive, or hateful content	Generating images of Christian churches only when prompted to depict "a house of worship" (Qadri et al., 2023a)
Misinformation Harms	AI systems generating and facilitating the spread of inaccurate or misleading information that causes people to develop false beliefs	An AI-generated image that was widely circulated on Twitter led several news outlets to falsely report that an explosion had taken place at the US Pentagon, causing a brief drop in the US stock market (Alba, 2023)
Information & Safety Harms	AI systems leaking, reproducing, generating or inferring sensitive, private, or hazardous information	An AI system leaks private images from the training data (Carlini et al., 2023a)
Malicious Use	AI systems reducing the costs and facilitating activities of actors trying to cause harm (e.g. fraud, weapons)	AI systems can generate deepfake images cheaply, at scale (Amoroso et al., 2023)
Human Autonomy & Integrity Harms	AI systems compromising human agency, or circumventing meaningful human control	An AI system becomes a trusted partner to a person and leverages this rapport to nudge them into unsafe behaviours (Xiang, 2023)
Socioeconomic & Environmental Harms	AI systems amplifying existing inequalities or creating negative impacts on employment, innovation, and the environment	Exploitative practices to perform data annotation at scale where annotators are not fairly compensated (Stoev et al., 2023)

Weidinger, Laura, et al. "Sociotechnical safety evaluation of generative AI systems."

Putting Ethics into Action: Responsible AI and Science

- For research, reproducibility is paramount
- But sometimes, conflict with privacy, especially with sensitive subjects
- Also conflicts with dual use — malicious usage
- Recommendations
  - Share your code
  - Share your data
- If using human subjects, e.g., for interviews or annotations:
  - Pay them fairly (especially if they are not co-authors)
  - Report their demographics (without compromising privacy!) — anonymization

NeurIPS 2021 Paper Checklist Guidelines

The NeurIPS Paper Checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. For each question in the checklist:

- You should answer yes, no, or n/a.
- You should reference the section(s) of the paper that provide support for your answer.
- You can also optionally write a short (1–2 sentence) justification of your answer.

You are encouraged to take a look at the checklist early on as it likely will ...

Reviewers will be asked to use the checklist as one of the factors in their e... ...ency: "Did you include [information]?" While "yes" is generally preferable to "no", it is perfectly acceptable to answer "no" provided a prop... ...it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). **In general, answering "no" or "n/a" is not grounds for rejection.** While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgement and write a justification to elaborate. All supporting evidence can appear either in the **main paper or the supplemental material**.

We provide guidance on how to answer each question below. You may additionally refer to our blog post introducing the checklist to learn more about its motivation and how it was created.

1. For all authors...
   - (a) Do the **main clai**...
     - Claims in the pa...
     - The paper's con... long as it is clea...
   - (b) Have you read th...
     - Please read the...

https://neurips.cc/Conferences/2021/PaperInformation/PaperChecklist

'Just What do You Think You're Doing, Dave?'* A Checklist for Responsible Data Use in NLP

Anna Rogers
Center for Social Data Science
University of Copenhagen
Denmark
arogers@sodas.ku.dk

Timothy Baldwin
CIS
The University of Melbourne
Australia
tb@ldwin.net

Kobi Leins
King's College / London
via Australia
kobi.leins@kcl.ac.uk

Abstract

A key part of the NLP ethics movement is responsible use of data, but exactly what that means or how it can be best achieved remain unclear. This position paper discusses the core legal and ethical principles for collection and sharing of textual data, and the tensions between them. We propose a potential ... contribute to the development of a consistent standard for data (re-)use, embraced across

(Mulligan et al., 2019; Xiang and Raji, 2019) and "bias" (Blodgett et al., 2020). The result is differing expectations and standards, which makes it harder for fellow researchers to understand the terms of (re-)use of new resources and models, and also complicates peer review.

Experts in biomedical data research have concluded that "truly informed consent ... requires ...ormational ...nce ...hat ...hers" (Mittelstadt, 2019). The AI community has already started to work on the latter (Mantelero, 2018;

Rogers, Anna, Timothy Baldwin, and Kobi Leins. "Just What Do You Think You're Doing, Dave? A Checklist for Responsible Data Use in NLP." *Findings of the Association for Computational Linguistics: EMNLP 2021.*

Revisiting course outcomes: After this course, you would have:

- Learned LLM fundamentals and their applications:
  - Foundations of Large Language Models
  - The recipe for building LLMs
  - Hands-on practice with prompting and fine-tuning LLMs
- Learned about Computational Social Science:
  - Main methods used in CSS
  - Applications of LLMs to CSS research questions and applications
  - CSS approaches to studying LLMs
- Gotten experience with a hands-on CSS project that uses LLMs

Upcoming Courses

IS 662	Network Science ☀️
IS 557	Public Blockchains ☀️
IE/IS 661	Text Analytics ☃️
IS 628	Seminar Advances in Public Blockchain ☃️
CS 721	Seminar Data Science I (Methods) ☀️/☃️
IS 723	Seminar Data Science II (Empirical Studies) ☀️/☃️
	Team Projects ☀️/☃️
	Master's Theses ☀️/☃️

[up-to-date list online]

New course!
IS 618 Social Media Data Analysis ☀️

Announcements

- Final presentations: 04.12.25, 11-13:45 PM in room 314-315, L 15, 1–6, 3rd floor
  - See directions
  - Ring the bell and we will let you in

- Each team has 15-20 mins to present, followed by 10 minutes of Q&A from the other students and instructors. Do NOT exceed 20 minutes.
  - Remember there 20% of the points for this course is based on class participation. You can get some of these points by asking constructive and polite questions to the other teams
  - Check grading rubric for what should be in the final presentations

- Please upload your final slides on ILIAS by 4.12.25 10:30 AM

- Final reports are due on 08.01.26 23:59 CET — upload on ILIAS

Questions?